

Decision-theoretic Planning

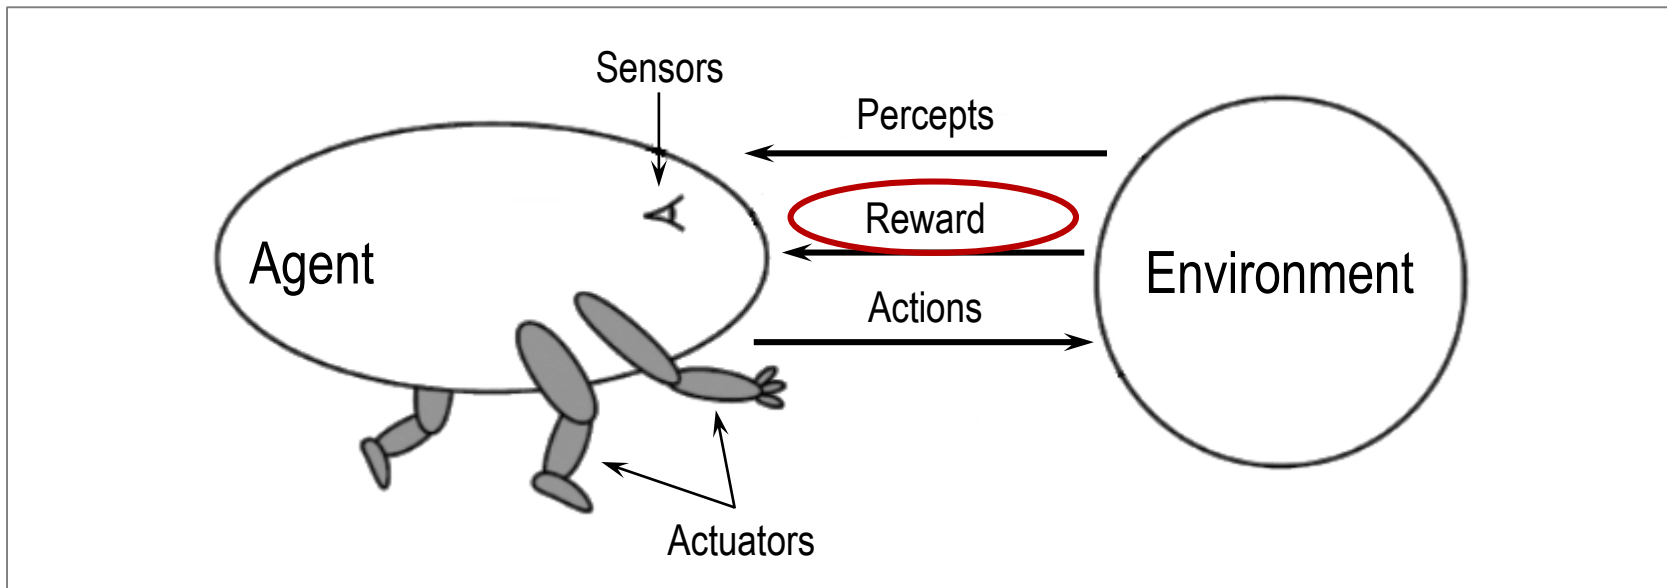


School of Electronic and Computer Engineering
Peking University

Wang Wenmin

What is Decision-theoretic Planning 什么是决策理论规划

- ❑ Classic planning is to find a plan to achieve its goals with lowest cost.
经典规划是寻找一个以最小代价到达其目标的计划。
- ❑ Decision-theoretic Planning is to find a plan to achieve its goals with maximum expected utility (MEU).
决策理论规划是寻找一个以最大期望效用 (MEU) 到达其目标的规划。



Example: Grid World 方格世界


- Agent lives in a grid, walls block agent's path. Stochastic movement.

智能体在格子中，围墙挡住了智能体的去路。随机移动。

- *Transition model*: 转换模型

- probability 0.8: agent moves up;
概率0.8: 智能体上移;
- probability 0.1: agent moves right or left;
概率0.1: 智能体左移、右移;
- no movement: if a wall in the direction;
不移: 若前方是堵墙;
- reward +1 and -1: two terminal states;
回报+1和-1: 两个终点状态;
- reward -0.04: other no-terminal states.
回报-0.04: 其它非终点状态。

- Goal: maximize sum of rewards. 目标: 回报值最大化。



How to Formulize and Solve 如何形式化与求解

□ How to formalize the problems of Decision-theoretic Planning?

如何对决策理论规划问题进行形式化？

■ Markov Decision Process (MDP)

马尔科夫决策过程 (MDP)

□ How to solve the problems of Markov Decision Process?

如何对马尔科夫决策过程进行求解？

■ Dynamic Programming

动态规划



Contents

- ☐ 8.5.1. Markov Decision Process
- ☐ 8.5.2. Dynamic Programming

Markov Decision Process (MDP) 马柯夫决策过程 (MDP)

- It is a *discrete time stochastic control process*, means action outcomes depend only on the current state.

是一种离散时间随机控制过程，意味着动作结果仅仅依赖于当前状态。

- A Markov Decision Process (MDP) is a 5-tuple (S, A, T, R, γ) , where

一个马柯夫决策过程是一个5元组 (S, A, P, R, γ) ，其中

- a set of **states**, $s \in S$ 一个状态集, $s \in S$
- a set **actions**, $a \in A$ 一个动作集, $a \in A$
- a **transition model**, $T(s, a, s')$ 一个迁移模型, $T(s, a, s')$
Probability that a from s leads to s' , i.e., $P(s' | s, a)$
从 s 导出 s' 的概率, 即: $P(s' | s, a)$
- a **reward function**, $R(s, a, s')$ 一个回报函数, $R(s, a, s')$
- discount, $\gamma \in [0, 1]$ 衰减, $\gamma \in [0, 1]$

Core Problem 核心问题

- The core problem of **classical planning**: 经典规划的核心问题
 - agent is in a **deterministic environment**,
智能体是在一个确定性的环境,
 - solving the problem is to find a **plan** to achieve its goal.
求解该问题是找到到一个达其目标的计划。
- The core problem of **Markov Decision Process (MDP)**: 马尔科夫决策过程的核心问题
 - agent is in a **discrete time stochastic environment**,
智能体处于一个离散时间随机环境,
 - solving the problem is to find a **policy** to control his process.
求解该问题是找到一个控制其过程的策略。

Finding policy is the core problem to solve MDPs

Core Problem 核心问题

- Given a MDP(S, A, T, R, γ), a policy is a computable function π that outputs for each state s an action a .

给定一个MDP(S, A, T, R, γ), 一个策略是一个计算函数 π , 它对每个状态 s 生成一个动作 a .

- A *deterministic policy* π is defined as: 一个确定性策略被定义为:

$$\pi : S \rightarrow A$$

- A *stochastic policy* π can also be defined as: 一个随机策略也可以被定义为:

$$\pi : S \times A \rightarrow [0, 1]$$

where $\pi(s, a) \geq 0$ and $\sum_a \pi(s, a) = 1$

- Goal is to choose a policy π that will maximize some cumulative function of the random rewards.

目标是选择一个策略 π , 使随机回报值的一些累积函数最大化。

Utilities and Optimal Policies 效用和优化策略

- In sequential decision problems, preferences are expressed between sequences of states.

在顺序决策问题中，偏好由状态顺序之间的顺序来表示。

- Usually use an additive utility functions:

通常采用一个累加效用函数：

$$U([s_0, s_1, s_2, \dots]) = R(s_0) + R(s_1) + R(s_2) + \dots = \sum_i R(s_i)$$

- Utility of a *state* (a.k.a. its value) is defined to be:

一个状态（亦称其值）的效用被定义为：

$$U(s_i) = \text{expected sum of rewards until termination assuming optimal actions.}$$

假设最佳动作结束之前的预期回报值的总和

- Two optimal policies: **Value Iteration** and **Policy Iteration**.

两个优化策略：值迭代和策略迭代。

1) Value Iteration 值迭代

□ Basic idea: 基本思想

- calculate the utility of each state, and then use the state utilities to select an optimal action in each state.

计算每个状态的效用，然后使用该状态效用在每个状态中选择一个最佳动作。

- π function is not used; instead the **value of π** is calculated within $U(s)$.

不使用 π 函数；而 π 值在 $U(s)$ 中计算。

□ Bellman equation for utilities:

贝尔曼效用等式：

$$U(s) = R(s) + \gamma \max_{\alpha \in A(s)} \sum_{s'} P(s' | s, \alpha) U(s')$$

Bellman equation is the basis of value iteration algorithm.

1) Value Iteration 值迭代

```

function VALUE-ITERATION( $mdp, \epsilon$ ) returns a utility function
  inputs:  $mdp$ , an MDP with states  $S$ , actions  $A(s)$ , transition model  $P(s' | s, a)$ ,
           rewards  $R(s)$ , discount  $\gamma$ 
            $\epsilon$ , the maximum error allowed in the utility of any state
  local variables:  $U, U'$ , vectors of utilities for states in  $S$ , initially zero
                     $\delta$ , the maximum change in the utility of any state in an iteration

  repeat
     $U \leftarrow U'; \delta \leftarrow 0$ 
    for each state  $s$  in  $S$  do
       $U'[s] \leftarrow R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$ 
      if  $|U'[s] - U[s]| > \delta$  then  $\delta \leftarrow |U'[s] - U[s]|$ 
  until  $\delta < \epsilon(1 - \gamma)/\gamma$ 
  return  $U$ 

```

The value iteration algorithm for calculating utilities of states.

计算状态效用的值迭代算法

2) Policy Iteration 策略迭代

□ Basic idea: alternate the two phases.

基本思想：交替执行如下两个阶段：

■ Policy evaluation: 策略迭代

given a **policy** π_i , calculate utility U_i of each state if π_i were to be executed.

给定一个策略 π_i ，如果 π_i 被执行的话，计算每个状态的效用 U_i 。

$$U_i(s) = R(s) + \gamma \sum_{s'} P(s' | s, \pi_i(s)) U_i(s')$$

■ Policy improvement: 策略改善

calculate a new MEU (maximum expected utility) policy π_{i+1} , using one-step look-ahead based on U_i .

使用基于 U_i 的提前看一步法，计算一个新的MEU（最大期待效用）策略 π_{i+1} 。

$$\pi^*(s) = \gamma \operatorname{argmax}_{\alpha \in A(s)} \sum_{s'} P(s' | s, \alpha) U(s')$$

2) Policy Iteration 策略迭代

```

function POLICY-ITERATION(mdp) returns a policy
  inputs: mdp, an MDP with states  $S$ , actions  $A(s)$ , transition model  $P(s' | s, a)$ 
  local variables:  $U$ , a vector of utilities for states in  $S$ , initially zero
                    $\pi$ , a policy vector indexed by state, initially random

  repeat
     $U \leftarrow \text{POLICY-EVALUATION}(\pi, U, \text{mdp})$ 
    unchanged?  $\leftarrow$  true
    for each state  $s$  in  $S$  do
      if  $\max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s'] > \sum_{s'} P(s' | s, \pi[s]) U[s']$  then do
         $\pi[s] \leftarrow \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$ 
        unchanged?  $\leftarrow$  false
  until unchanged?
  return  $\pi$ 

```

The policy iteration algorithm for calculating an optimal policy.

计算最佳策略的值迭代算法

Thank you for your attention!

AI