

人工智能技术与应用

搜索II

2019.03.12

大纲

- 搜索：路径无关
- 约束满足问题CSP
- 约束优化问题COP
- AI在金融市场应用

建模-推理-学习

- 模型
 - 状态空间模型
- 推理
 - 最短路算法
- 学习
 - 感知机学习

Modeling

Inference

Learning

Search where the path doesn't matter

- So far, looked at problems where the path was the solution
 - Traveling on a graph
 - Eights puzzle
- However, in many problems, we just want to find a goal state
 - Doesn't matter how we get there


Queens puzzle

- Place eight queens on a chessboard so that no two attack each other


				Q			
		Q					
Q							
						Q	
	Q						
							Q
					Q		
			Q				

Search formulation of the queens puzzle

- **Successors:** all valid ways of placing additional queen on the board; **goal:** eight queens placed




How big is this tree?
How many leaves?
What if they were rooks?



Search formulation of the queens puzzle

- **Successors:** all valid ways of placing a queen in the next column; **goal:** eight queens placed




Search tree size?
What if they were rooks?
What kind of search is best?

Constraint satisfaction problems (CSPs)

- Defined by:
 - A set of **variables** x_1, x_2, \dots, x_n
 - A **domain** D_i for each variable x_i
 - **Constraints** c_1, c_2, \dots, c_m
- A constraint is specified by
 - A subset (often, two) of the variables
 - All the allowable joint assignments to those variables
- Goal: find a **complete, consistent** assignment
- Queens problem: (other examples in next slides)
 - x_i in $\{1, \dots, 8\}$ indicates in which row in the i th column to place a queen
 - For example, constraint on x_1 and x_2 : $\{(1,3), (1,4), (1,5), (1,6), (1,7), (1,8), (2,4), (2,5), \dots, (3,1), (3,5), \dots \dots\}$

Graph coloring

- Fixed number of colors; no two adjacent nodes can share a color



Satisfiability

- Formula in conjunctive normal form:

$(x_1 \text{ OR } x_2 \text{ OR NOT}(x_4)) \text{ AND } (\text{NOT}(x_2) \text{ OR NOT}(x_3)) \text{ AND } \dots$

- Label each variable x_j as true or false so that the formula becomes true

Constraint hypergraph:

each hyperedge

represents a constraint

Cryptarithmic puzzles

$$\begin{array}{r} \text{TWO} \\ \text{TWO} + \\ \hline \text{FOUR} \end{array}$$

E.g., setting $F = 1$, $O = 4$, $R = 8$, $T = 7$, $W = 3$, $U = 6$ gives $734 + 734 = 1468$

Cryptarithmic puzzles...


$$\begin{array}{r} \text{TWO} \\ \text{TWO} + \\ \hline \text{FOUR} \end{array}$$

Trick: introduce **auxiliary variables** X, Y

$$O + O = 10X + R$$

$$W + W + X = 10Y + U$$

$$T + T + Y = 10F + O$$



What would the search tree look like?


also need pairwise constraints
between original variables if they
are supposed to be different

Generic approaches to solving CSPs

- State: some variables assigned, others not assigned
- Naïve successors definition: any way of assigning a value to an unassigned variable results in a successor
 - Can check for consistency when expanding
 - How many leaves do we get in the worst case?
- CSPs satisfy **commutativity**: order in which actions applied does not matter
- Better idea: only consider assignments for a single variable at a time
 - How many leaves?

Choice of variable to branch on is still flexible!

- Do not always need to choose same variable at same level
- Each of variables A , B , C takes values in $\{0,1\}$



- Can you **prove** that this never increases the size of the tree?


A generic recursive search algorithm

(*assignment* is a partial assignment)

- **Search(*assignment*, *constraints*)**
- If *assignment* is complete, return it
- Choose an unassigned variable *x*
- For every value *v* in *x*'s domain, if setting *x* to *v* in *assignment* does not violate *constraints*:
 - Set *x* to *v* in *assignment*
 - *result* := Search(*assignment*, *constraints*)
 - If *result* != *failure* return *result*
 - Unassign *x* in *assignment*
- Return *failure*

Keeping track of remaining possible values

- For every variable, keep track of which values are still possible



only one possibility for
last column; might as
well fill in

now only one left for
other two columns

done!
(no real branching
needed!)

- General heuristic: branch on variable with fewest values remaining

Arc consistency

- Take two variables connected by a constraint
- Is it true that for **every** remaining value d of the first variable, there exists **some** value d' of the other variable so that the constraint is satisfied?
 - If so, we say the **arc from the first to the second variable is consistent**
 - If not, can remove the value d
- General concept: **constraint propagation**


Q				x			x
				x			
		Q		x			x
				x			
				x	Q		x
				x			

Consider cryptarithmic puzzle again...

Is the arc from the fifth to the eighth column consistent?

What about the arc from the eighth to the fifth?


An example where arc consistency fails



- $A = B, B = C, C \neq A$ – obviously inconsistent
 - \sim Moebius band
- However, arc consistency cannot eliminate anything

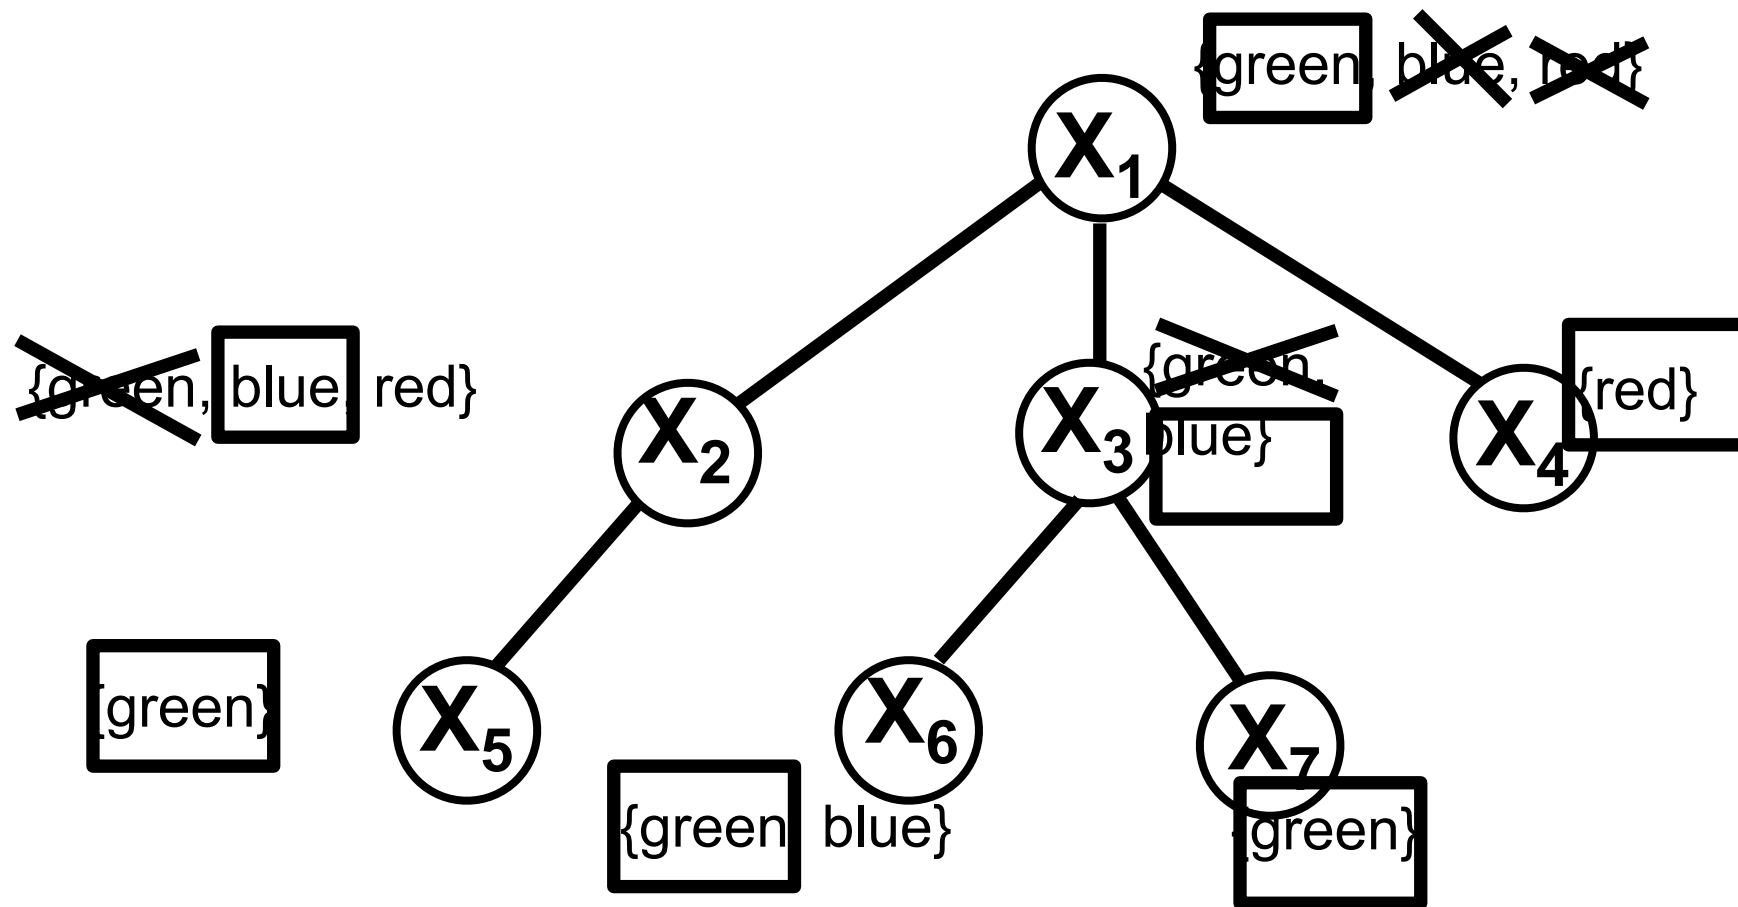
Tree-structured constraint graphs

- Suppose we only have pairwise constraints and the graph is a **tree** (or **forest** = multiple disjoint trees)



- Dynamic program for solving this (linear in #variables):
 - Starting from the leaves and going up, for each node x , compute all the values for x such that the subtree rooted at x can be solved
 - Equivalently: apply arc consistency from each parent to its children, starting from the bottom
 - If no domain becomes empty, once we reach the top, easy to fill in solution


Example: graph coloring with limited set of colors per node



- Stage 1: moving upward, cross out the values that cannot work with the subtree below that node
- Stage 2: if a value remains at the root, there is a solution: go downward to pick a solution

Generalizations of the tree-based approach

- What if our constraint graph is “almost” a tree?




- A **cycle cutset** is a set of variables whose removal results in a tree (or forest)
 - E.g. $\{X_1\}$, $\{X_6\}$, $\{X_2, X_3\}$, $\{X_2, X_4\}$, $\{X_3, X_4\}$
- Simple algorithm: for every internally consistent assignment to the cutset, solve the remaining tree as before (runtime?)
- Graphs of **bounded treewidth** can also be solved in polynomial time (won't define these here)

A different approach: optimization


- Let's say every way of placing 8 queens on a board, one per column, is **feasible**
- Now we introduce an **objective**: minimize the number of pairs of queens that attack each other
 - More generally, minimize the number of violated constraints
- Pure optimization

Local search: hill climbing

- Start with a complete state
- Move to successor with best (or at least better) objective value
 - Successor: move one queen within its column



- Local search can get stuck in a **local optimum**



Avoiding getting stuck with local search

- **Random restarts**: if your hill-climbing search fails (or returns a result that may not be optimal), restart at a random point in the search space
 - Not always easy to generate a random state
 - Will **eventually** succeed (why?)
- **Simulated annealing**:
 - Generate a random successor (possibly worse than current state)
 - Move to that successor with some probability that is sharply decreasing in the badness of the state
 - Also, over time, as the “temperature decreases,” probability of bad moves goes down


Constraint optimization

- Like a CSP, but with an objective
 - E.g., minimize number of violated constraints
 - Another example: no two queens can be in the same row or column (hard constraint), minimize number of pairs of queens attacking each other diagonally (objective)
- Can use all our techniques from before: heuristics, A^* , IDA^* , ...
- Also popular: **depth-first branch-and-bound**
 - Like depth-first search, except do not stop when first feasible solution found; keep track of best solution so far
 - Given admissible heuristic, do not need to explore nodes that are worse than best solution found so far


Minimize #violated diagonal constraints

- **Cost of a node:** #violated diagonal constraints so far

- No heuristic
*(matter of definition;
could just as well say
that violated constraints
so far is the heuristic
and interior nodes have
no cost)*




cost = 0




cost = 0




cost = 1




cost = 0



Depth first branch and bound will find a suboptimal solution here first (no way to tell at this point this is worse than right node)





A* (=uniform cost here), IDA*
(=iterative lengthening here) will
never explore this node



Optimal solution is down here
(cost 0)

Linear programs: example

- We make reproductions of two paintings



maximize $3x + 2y$

subject to

$$4x + 2y \leq 16$$

$$x + 2y \leq 8$$

$$x + y \leq 5$$

- Painting 1 sells for \$30, painting 2 sells for \$20
- Painting 1 requires 4 units of blue, 1 green, 1 red $x \geq 0$
- Painting 2 requires 2 blue, 2 green, 1 red $y \geq 0$
- We have 16 units blue, 8 green, 5 red

Solving the linear program graphically

maximize $3x + 2y$

subject to


$$4x + 2y \leq 16$$

$$x + 2y \leq 8$$

$$x + y \leq 5$$

$$x \geq 0$$

$$y \geq 0$$



Modified LP

maximize $3x + 2y$

subject to

$$4x + 2y \leq 15$$

$$x + 2y \leq 8$$

$$x + y \leq 5$$

$$x \geq 0$$

$$y \geq 0$$

Optimal solution: $x = 2.5$, $y = 2.5$

Solution value = $7.5 + 5 = 12.5$

Half paintings?

Integer (linear) program

maximize $3x + 2y$

subject to


$$4x + 2y \leq 15$$

$$x + 2y \leq 8$$

$$x + y \leq 5$$

$x \geq 0$, integer

$y \geq 0$, integer



Mixed integer (linear) program

maximize $3x + 2y$

subject to


$$4x + 2y \leq 15$$

$$x + 2y \leq 8$$

$$x + y \leq 5$$

$$x \geq 0$$


$$y \geq 0, \text{ integer}$$



Solving linear/integer programs

- Linear programs can be solved efficiently
 - Simplex, ellipsoid, interior point methods...
- (Mixed) integer programs are NP-hard to solve
 - Quite easy to model many standard NP-complete problems as integer programs (try it!)
 - Search type algorithms such as branch and bound
- Standard packages for solving these
 - GNU Linear Programming Kit, CPLEX, ...
- **LP relaxation** of (M)IP: remove integrality constraints
 - Gives upper bound on MIP (~admissible heuristic)

Graph coloring as an integer program



- Let's say $x_{B,green}$ is 1 if B is colored green, 0 otherwise
- Must have $0 \leq x_{B,green} \leq 1$, $x_{B,green}$ integer
 - shorthand: $x_{B,green} \in \{0,1\}$
- Constraint that B and C can't both be green: $x_{B,green} + x_{C,green} \leq 1$
- Etc.
- Solving integer programs is at least as hard as graph coloring, hence NP-hard (we have **reduced** graph coloring to IP)

Satisfiability as an integer program

$(x_1 \text{ OR } x_2 \text{ OR NOT}(x_4)) \text{ AND } (\text{NOT}(x_2) \text{ OR NOT}(x_3)) \text{ AND } \dots$

becomes

for all x_j , $0 \leq x_j \leq 1$, x_j integer (shorthand: x_j in $\{0,1\}$)

$$x_1 + x_2 + (1-x_4) \geq 1$$

$$(1-x_2) + (1-x_3) \geq 1$$

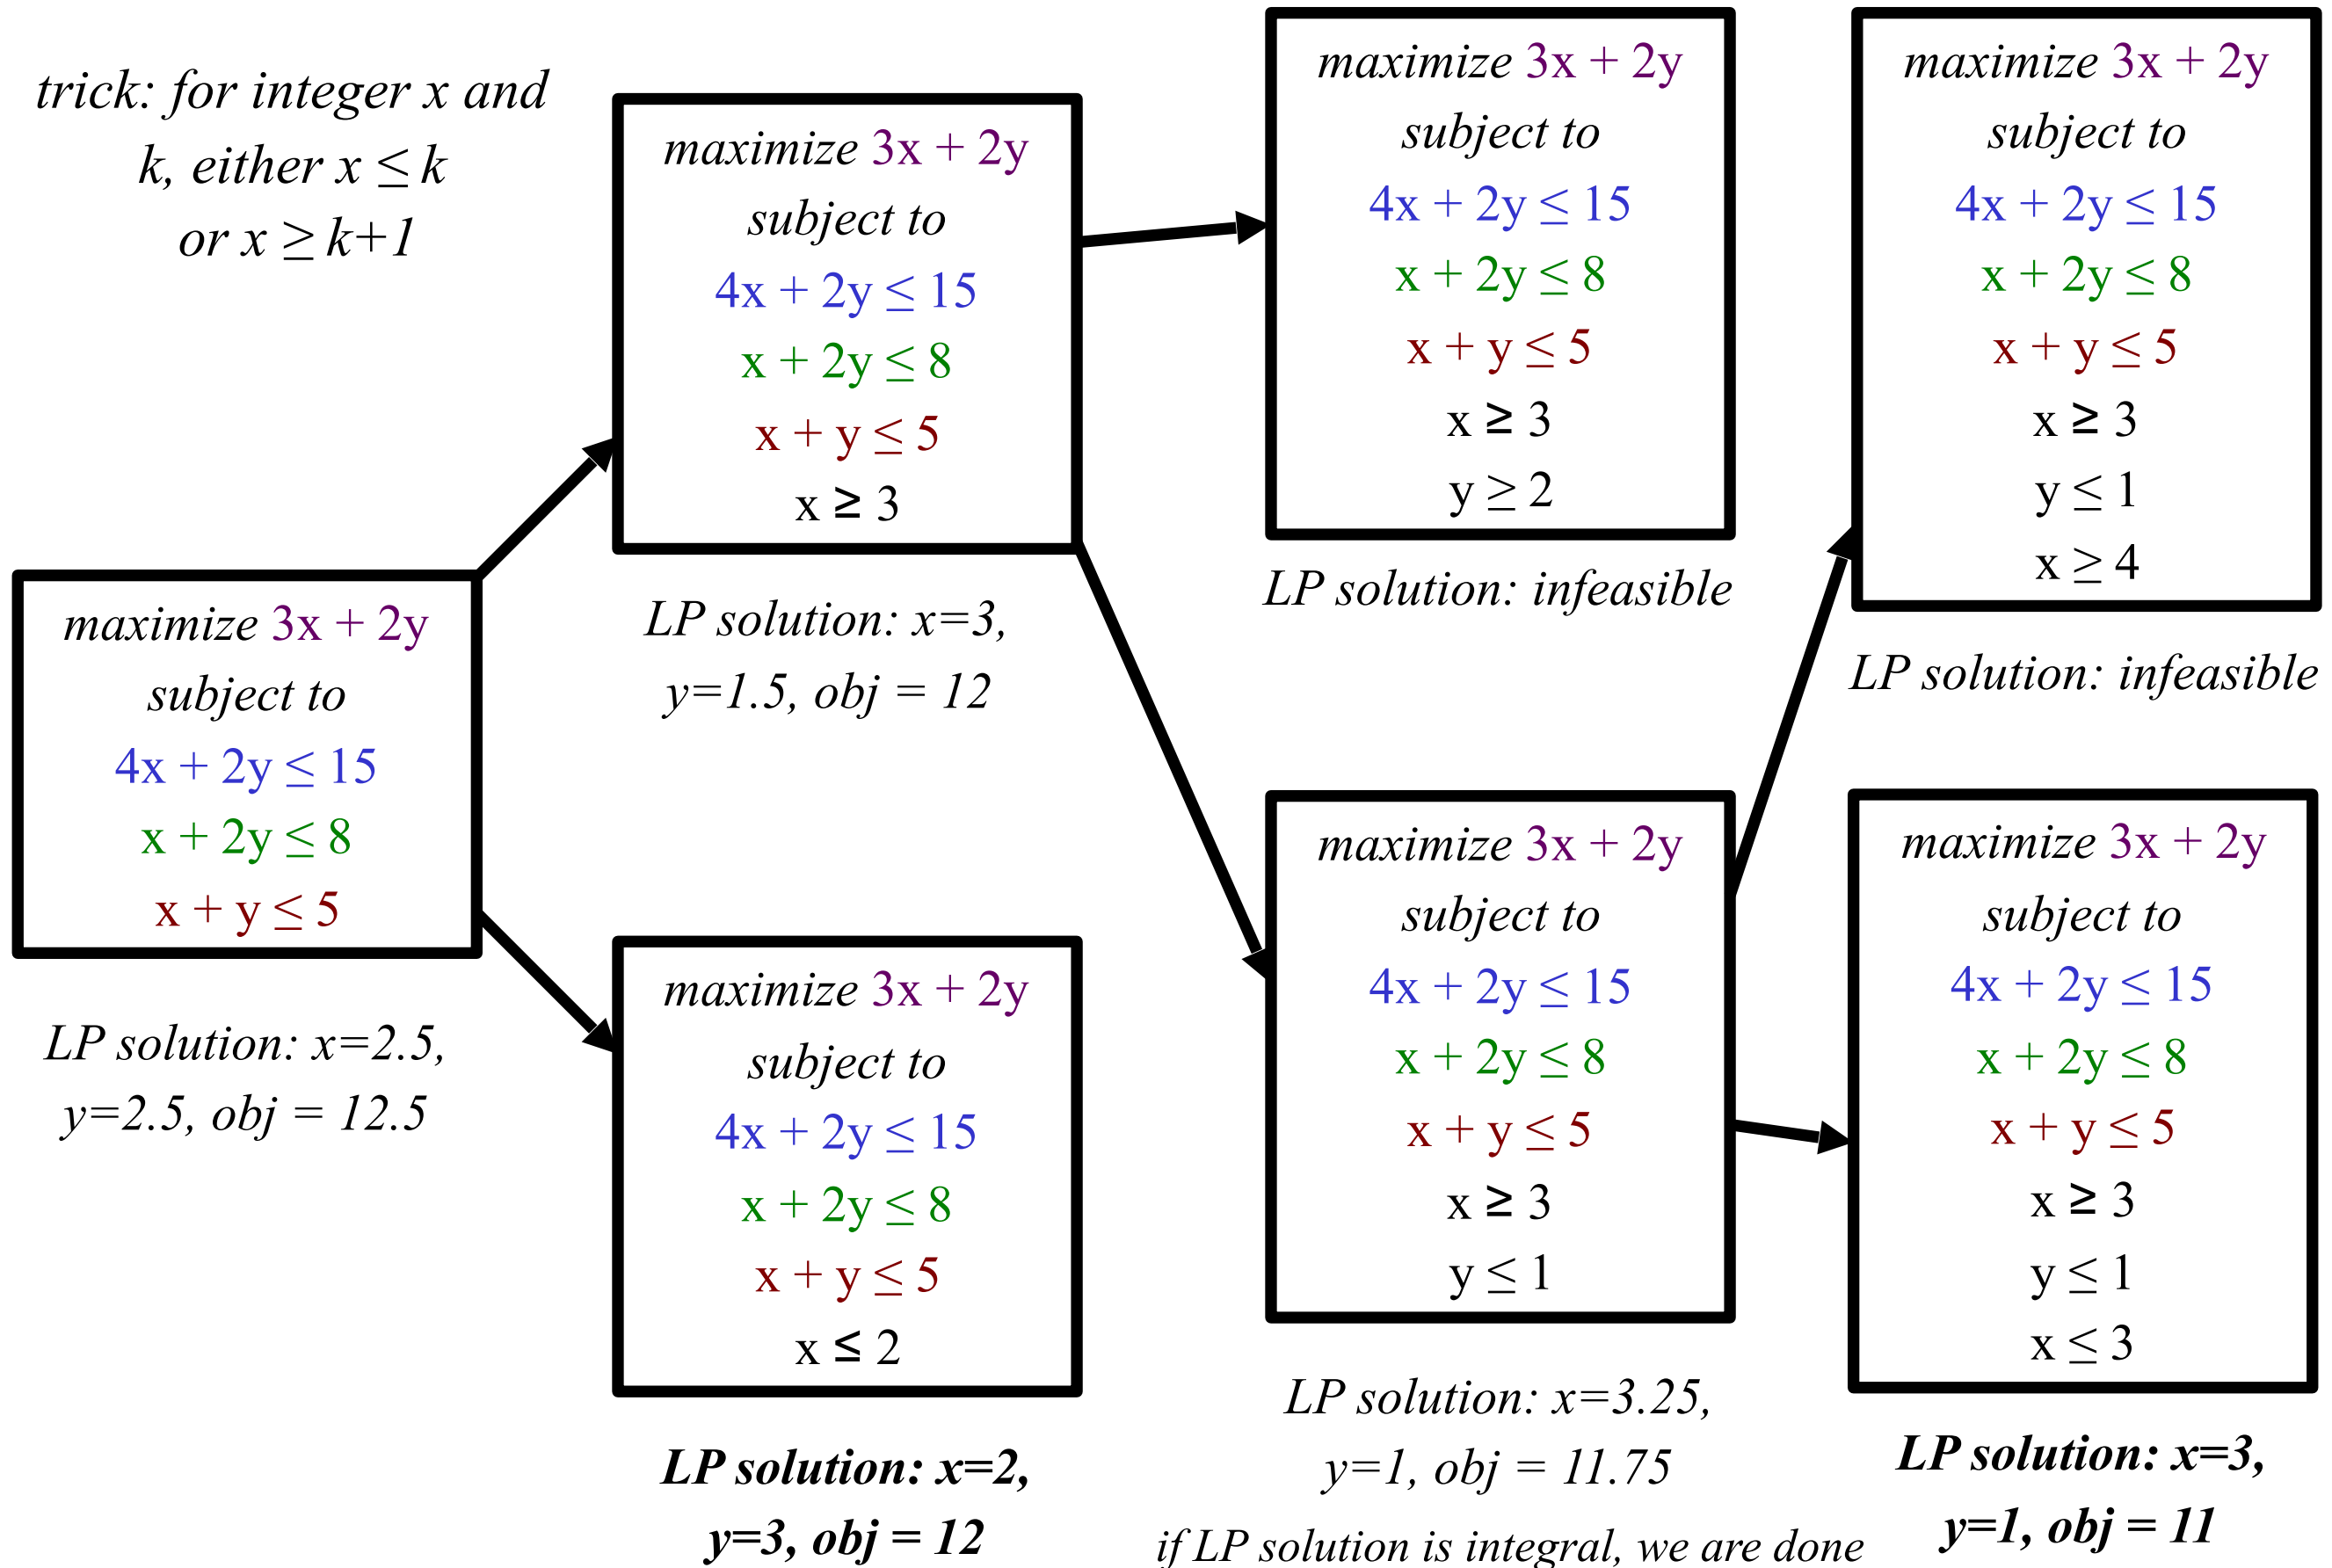
...

Solving integer programs is at least as hard as satisfiability, hence NP-hard (we have **reduced** SAT to IP)

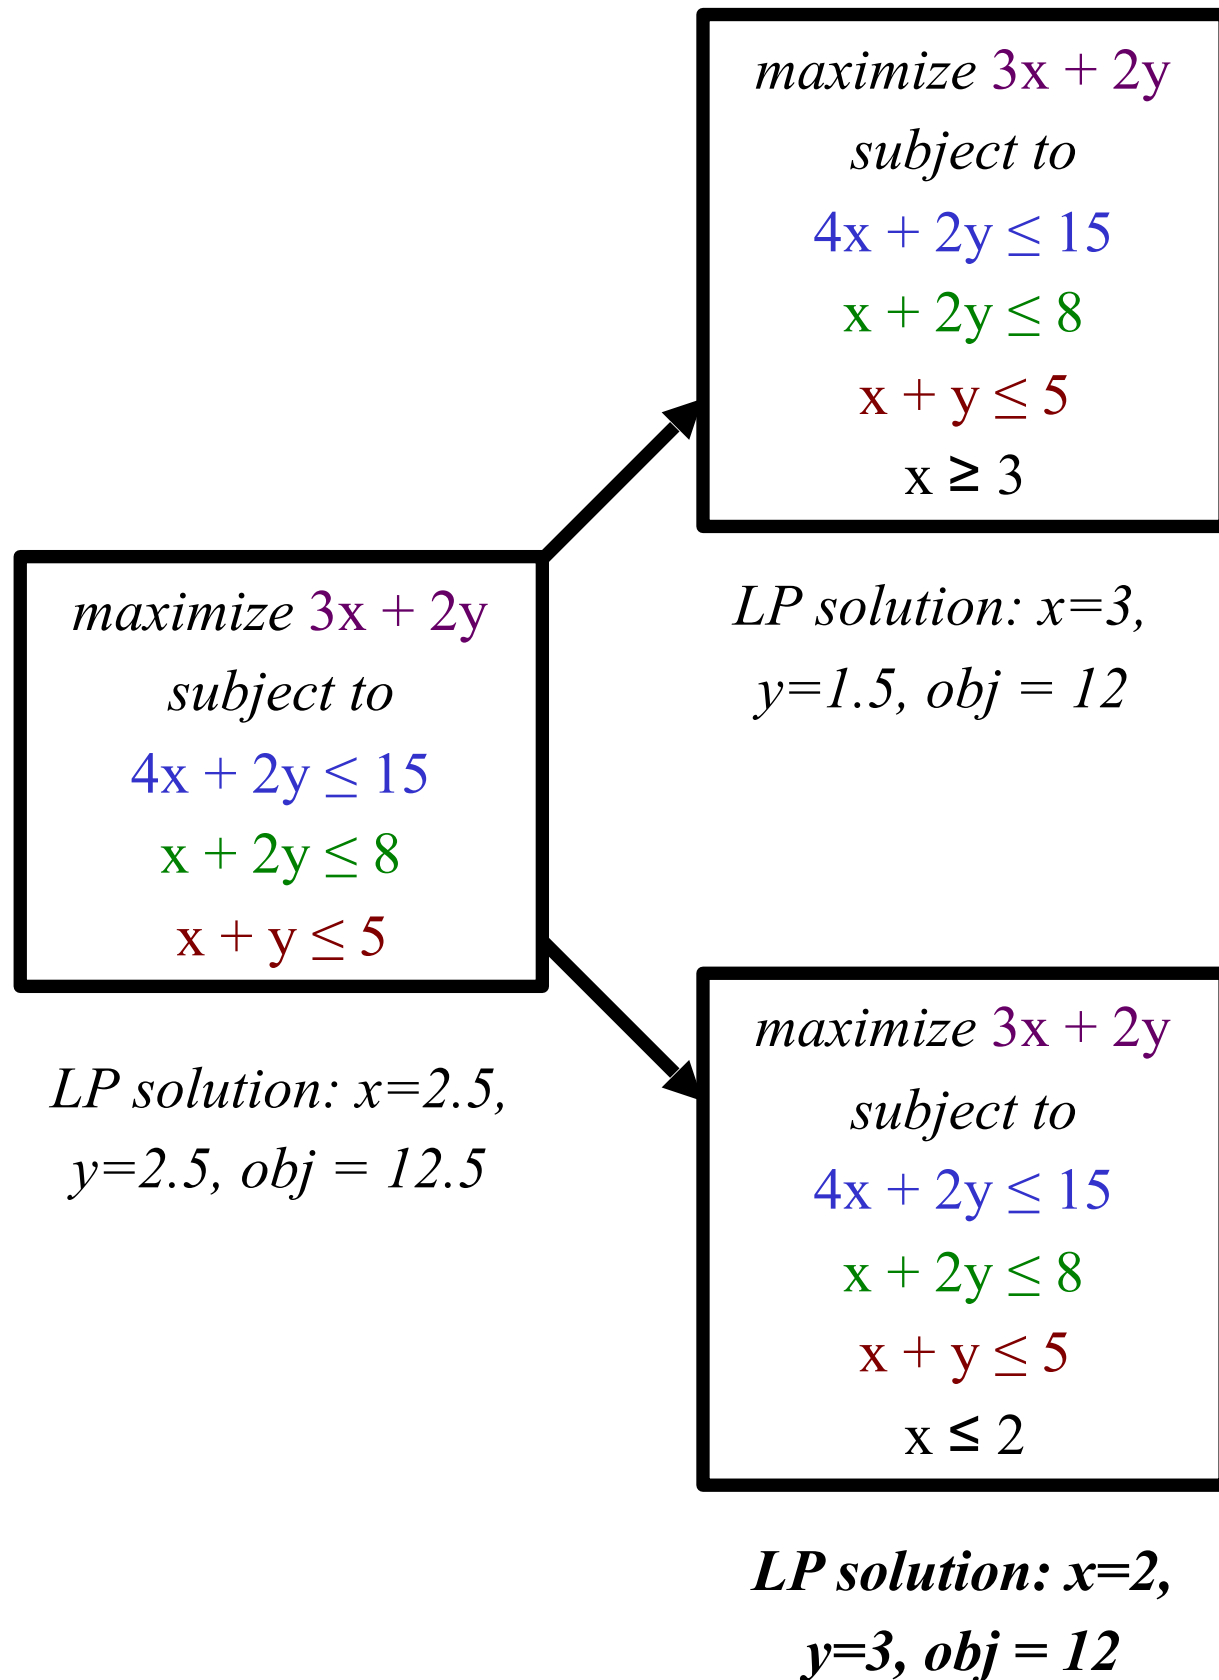
Try modeling other NP-hard problems as (M)IP!

Solving the integer program with DFS branch and bound

trick: for integer x and k , either $x \leq k$ or $x \geq k+1$



Again with a more fortunate choice




done!

AI+金融市场


看涨期权 (call option)

- A (European) call option $C(S, k, t)$ gives you the right to buy stock S at (strike) price k on (expiry) date t
 - American call option can be exercised early
 - European one easier to analyze
- How much is a call option worth at time t (as a function of the price of the stock)?




看跌期权 (put option)

- A (European) put option $P(S, k, t)$ gives you the right to sell stock S at (strike) price k on (expiry) date t
- How much is a put option worth at time t (as a function of the price of the stock)?




债券 bonds

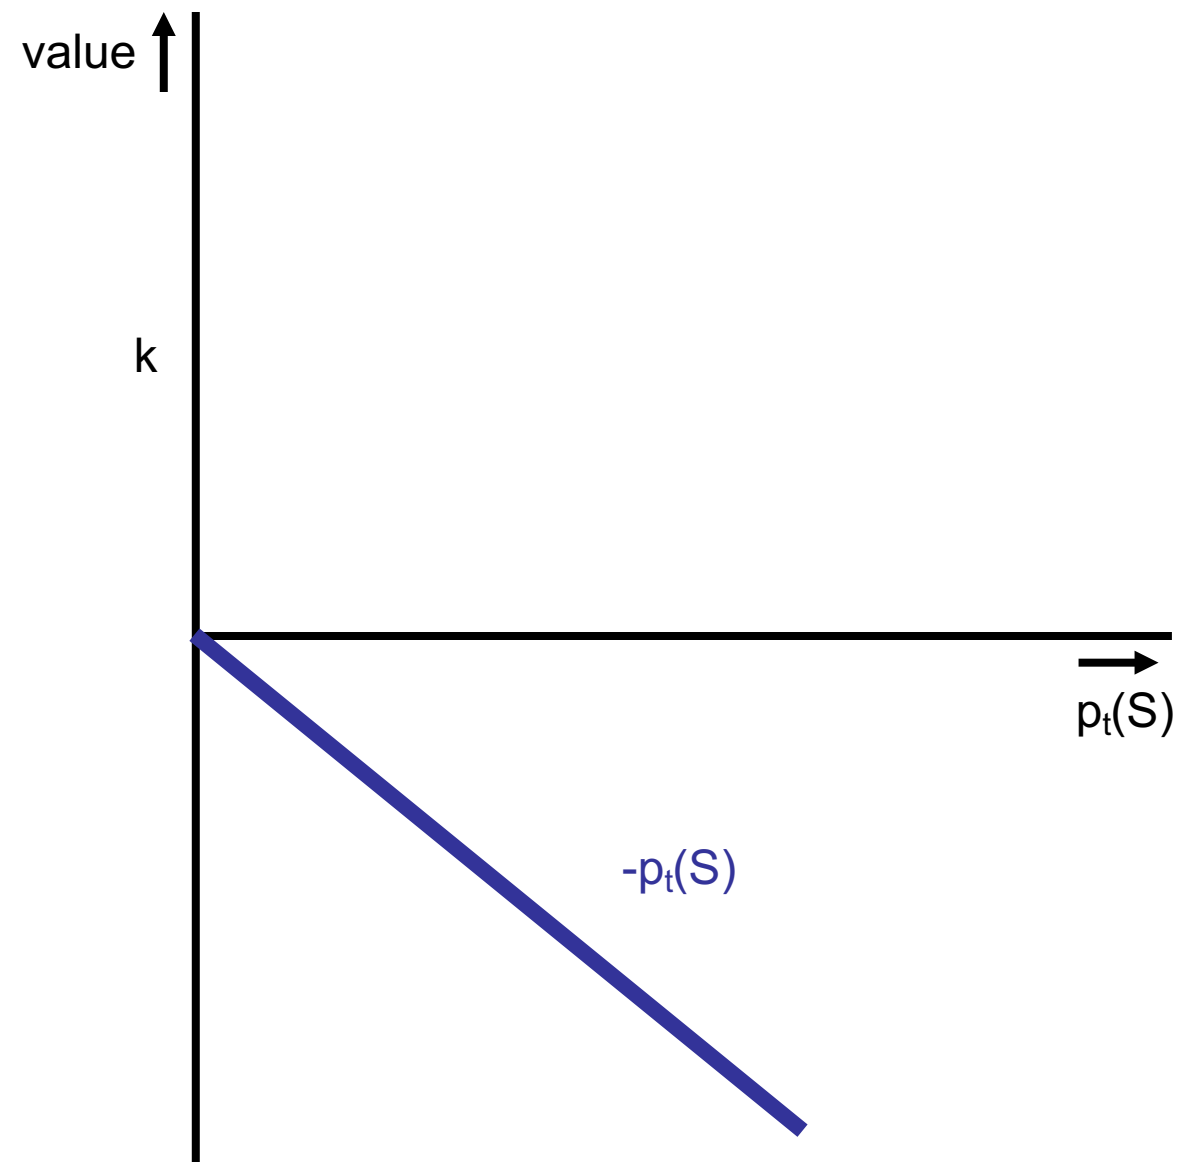
- A bond $B(k, t)$ pays off k at time t



股票 Stocks





Selling a stock (short)





A portfolio

- One call option $C(S, k, t)$ + one bond $B(k, t)$



Another portfolio

- One put option $P(S, k, t)$ + one stock S



Put-call parity (等价理论)


- $C(S, k, t) + B(k, t)$ will have the same value at time t as $P(S, k, t) + S$ (regardless of the value of S)
- Assume stocks pay no dividends
- Then, portfolio should have the same value at any time before t as well
- I.e., for any $t' < t$, it should be that
$$p_{t'}(C(S, k, t)) + p_{t'}(B(k, t)) = p_{t'}(P(S, k, t)) + p_{t'}(S)$$
- Arbitrage argument: suppose (say) $p_{t'}(C(S, k, t)) + p_{t'}(B(k, t)) < p_{t'}(P(S, k, t)) + p_{t'}(S)$
- Then: buy $C(S, k, t) + B(k, t)$, sell (short) $P(S, k, t) + S$
- Value of portfolio at time t is 0
- Guaranteed profit!

Another perspective: auctioneer

- **Auctioneer** receives buy and sell offers, has to choose which to accept
- E.g.: offers received: buy(S, \$10); sell(S, \$9)
- Auctioneer can accept both offers, profit of \$1
- E.g. (put-call parity):
 - sell(**C**(S, k, t), \$3)
 - sell(**B**(k, t), \$4)
 - buy(**P**(S, k, t), \$5)
 - buy(**S**, \$4)
- Can accept all offers at no risk!


“Butterfly” portfolio

- 1 call at strike price $k-c$
- -2 calls at strike k
- 1 call at strike $k+c$




Another portfolio

- Can we create this portfolio?



Yet another portfolio

- How about this one?



Securities conditioned on finite set of outcomes

- E.g., InTrade: security that pays off 1 if Trump is the Republican nominee in 2016
- Can we construct a portfolio that pays off 1 if Clinton is the Democratic nominee AND Trump is the Republican nominee?

	Trump not nom.	Trump nom.
Clinton not nom.	\$0	\$0
Clinton nom.	\$0	\$1









Arrow-Debreu securities

- Suppose S is the set of **all** states that the world can be in tomorrow
- For each s in S , there is a corresponding Arrow-Debreu security that pays off 1 if s happens, 0 otherwise
- E.g., s could be: Clinton is nominee and Trump is nominee and S_1 is at \$4 and S_2 at \$5 and butterfly 432123 flaps its wings in Peru and...
- Not practical, but conceptually useful
- Can think about Arrow-Debreu securities **within a domain** (e.g., states only involve stock trading prices)
- Practical for small number of states




With Arrow-Debreu securities you can do anything...

- Suppose you want to receive \$6 in state 1, \$8 in state 2, \$25 in state 3
- ... simply buy 6 AD securities for state 1, 8 for state 2, 25 for state 3
- Linear algebra: Arrow-Debreu securities are a basis for the space of all possible securities

The auctioneer problem

- Tomorrow there must be one of   
- Agent 1 offers \$5 for a security that pays off \$10 if  
- Agent 2 offers \$8 for a security that pays off \$10 if  
- Agent 3 offers \$6 for a security that pays off \$10 if 
- Can we accept some of these at offers **at no risk?**

Reducing auctioneer problem to ~combinatorial exchange winner determination problem

- Let (x, y, z) denote payout under respectively   
- Previous problem's bids:
 - 5 for $(0, 10, 10)$
 - 8 for $(10, 0, 10)$
 - 6 for $(10, 0, 0)$
- Equivalently:
 - $(-5, 5, 5)$
 - $(2, -8, 2)$
 - $(4, -6, -6)$
- Sum of accepted bids should be $(\leq 0, \leq 0, \leq 0)$ to have no risk
- Sometimes possible to partially accept bids

A bigger instance (4 states)

- Objective: maximize our **worst-case** profit
 - 3 for $(0, 0, 11, 0)$
 - 4 for $(0, 2, 0, 8)$
 - 5 for $(9, 9, 0, 0)$
 - 3 for $(6, 0, 0, 6)$
 - 1 for $(0, 0, 0, 10)$
-
- What if they are partially acceptable?

Settings with large state spaces

- Large = exponentially large
 - Too many to write down
- Examples:
- $S = S_1 \times S_2 \times \dots \times S_n$
 - E.g., $S_1 = \{\text{Clinton not nom.}, \text{Clinton nom.}\}$, $S_2 = \{\text{Trump not nom.}, \text{Trump nom.}\}$, $S = \{(-C, -T), (-C, +T), (+C, -T), (+C, +T)\}$
 - If all S_i have the same size k , there are k^n different states
- S is the set of all rankings of n candidates
 - E.g., outcomes of a horse race
 - $n!$ different states (assuming no ties)

Bidding languages

- How should **trader** (bidder) express preferences?
- Logical bidding languages [Fortnow et al. 2004]:
 - (1) “If Trump nominated OR (Cruz nominated AND Clinton nominated), I want to receive \$10; I’m willing to pay \$6 for this.”
- If the state is a ranking [Chen et al. 2007]:
 - (2a) “If horse A ranks 2nd, 3rd, or 4th I want to receive \$10; I’m willing to pay \$6 for this.”
 - (2b) “If one of horses A, C, D rank 2nd, I want to receive \$10; I’m willing to pay \$6 for this.”
 - (2c) “If horse A ranks ahead of horse C, I want to receive \$10; I’m willing to pay \$6 for this.”
- Winner determination problem is NP-hard for all of these, except for (2a) and (2b) which are in P **if** bids can be partially accepted

谢谢！